

```
In [1]: %matplotlib inline
```

```
In [2]: import tensorflow as tf
import tensorflow.contrib.slim as slim
from scipy.io import loadmat
import matplotlib
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import numpy as np
import os
import itertools
import math
from ipywidgets import interact, interactive, fixed
import ipywidgets as widgets
```

```
In [3]: import classifier_utils as utils
```

## Get the training data

```
In [4]: categories = [
        'epithelial',
        'fibroblast',
        'inflammatory',
        'others',
    ]
```

```
In [5]: # Read in the raw data
(raw_imgs, raw_centres, raw_labels) = utils.get_dataset(100, categories)
```

```
In [6]: # Extract example patches from the data
H = 27
W = 27
(patches, labels, centres, img_ids) = utils.get_examples(
    raw_imgs, raw_centres, raw_labels, H, W)
```

Dropped 2082 patches because too close to image border

```
In [7]: # Organize examples into training and test data
N = patches.shape[0]
num_train = int(0.8 * N)

np.random.seed(0) # predictable shuffling for now
perm = np.random.permutation(N)

train_patches, test_patches = np.split(patches[perm], [num_train])
train_labels, test_labels = np.split(labels[perm], [num_train])
train_centres, test_centres = np.split(centres[perm], [num_train])
train_img_ids, test_img_ids = np.split(img_ids[perm], [num_train])

# Convert to float
train_patches = train_patches / 255.0
test_patches = test_patches / 255.0
```

```
In [8]: # Sanity check
train_patches.shape, train_labels.shape, test_patches.shape,
test_labels.shape
```

```
Out[8]: ((16289, 27, 27, 3), (16289, 4), (4073, 27, 27, 3), (4073, 4))
```

```
In [9]: # Expand/augment the training data
desired_cnt_per_category = 15000
sorted_train_dict = utils.expand_training_data(
    raw_imgs, train_patches, train_labels, train_centres, train_img_ids,
    desired_cnt_per_category)
# Convert to float
sorted_train_dict['patches'] = sorted_train_dict['patches'] / 255.0
```

```
In [10]: # Suffle the augmented training data
trainN = sorted_train_dict['patches'].shape[0]

np.random.seed(123) # predictable shuffling for now
perm = np.random.permutation(trainN)
train_dict = {k : v[perm] for (k, v) in sorted_train_dict.iteritems()}
```

```
In [11]: sess = tf.InteractiveSession()
```

```
In [12]: def upscale(images, scales, name):  
        (yscale, xscale) = scales  
        (batch_size, height, width, channels) = images.get_shape().as_list()  
        with tf.name_scope(name):  
            return tf.image.resize_images(images, height*yscale,  
width*xscale,  
                                           method=tf.image.ResizeMethod.NEAREST  
T_NEIGHBOR)
```

```

def autoencoder_model_1(image_batch):
    with slim.arg_scope([slim.conv2d, slim.fully_connected],
                        activation_fn=tf.nn.relu,
                        weights_initializer=tf.truncated_normal_initializer(stddev=0.01),
                        biases_initializer=tf.zeros_initializer,
                        #weights_regularizer=slim.l2_regularizer(0.0005),
                        biases_regularizer=None):
        with slim.arg_scope([slim.conv2d, slim.conv2d_transpose],
                            padding='VALID'):
            with slim.arg_scope([slim.dropout], keep_prob=0.8):
                with slim.arg_scope([slim.conv2d, slim.fully_connected, slim.conv2d_transpose],
                                    normalizer_fn=slim.batch_norm):
                    net = image_batch

                    estack = []
                    def r(x):
                        estack.append(x)
                        return x
                    with tf.variable_scope("encoder"):
                        net = r(slim.conv2d(net, 36, [4, 4], scope='1_conv'))
                        net = r(slim.conv2d(net, 42, [3, 3], scope='2_conv'))
                        net = r(slim.max_pool2d(net, [2, 2], scope='3_max_pool'))
                        net = r(slim.conv2d(net, 48, [4, 4], scope='4_conv'))
                        net = r(slim.max_pool2d(net, [2, 2], scope='5_max_pool'))
                        net = r(slim.flatten(net, scope='5_flatten'))
                        net = r(slim.fully_connected(net, 256, scope='6_fc'))
                        #net = r(slim.dropout(net, scope='6_dropout'))
                        net = r(slim.fully_connected(net, 64, scope='7_fc'))

                    encoded = net

                    dstack = []
                    def r(x):
                        dstack.append(x)
                        return x
                    with tf.variable_scope("decoder"):
                        net = r(slim.fully_connected(net, 256, scope='1_fc'))
                        #net = r(slim.dropout(net, scope='1_dropout'))
                        net = r(slim.fully_connected(net, 4*4*48, scope='2_fc'))
                        net = r(tf.reshape(net, [-1, 4, 4, 48], name='2_reshape'))
                        net = r(upscale(net, [2, 2], name='3_upscale'))
                        net = r(slim.conv2d_transpose(net, 42, [4, 4], scope='4_conv'))
                        net = r(upscale(net, [2, 2], name='5_upscale'))
                        net = r(slim.conv2d_transpose(net, 36, [3, 3], scope='6_conv'))

```

```
ope='6_conv'))  
net = r(slim.conv2d_transpose(net, 3, [4, 4], scope='7_conv'))  
  
reconstructed = net  
  
return (encoded, reconstructed, estack, dstack)
```

```

def autoencoder_model_2(image_batch):
    with slim.arg_scope([slim.conv2d, slim.fully_connected],
                        activation_fn=tf.nn.relu,
                        weights_initializer=tf.truncated_normal_initializer(stddev=0.01),
                        biases_initializer=tf.zeros_initializer,
                        #weights_regularizer=slim.l2_regularizer(0.0005),
                        biases_regularizer=None):
        with slim.arg_scope([slim.conv2d, slim.conv2d_transpose],
                            padding='VALID'):
            with slim.arg_scope([slim.dropout], keep_prob=0.8):
                with slim.arg_scope([slim.conv2d, slim.fully_connected, slim.conv2d_transpose],
                                    normalizer_fn=slim.batch_norm):
                    net = image_batch

                    estack = []
                    def r(x):
                        estack.append(x)
                        return x
                    with tf.variable_scope("encoder"):
                        net = r(slim.conv2d(net, 36, [4, 4], scope='1_conv'))
                        net = r(slim.conv2d(net, 42, [3, 3], scope='2_conv'))
                        net = r(slim.max_pool2d(net, [2, 2], scope='3_max_pool'))
                        net = r(slim.conv2d(net, 48, [4, 4], scope='4_conv'))
                        net = r(slim.max_pool2d(net, [2, 2], scope='5_max_pool'))
                        net = r(slim.flatten(net, scope='5_flatten'))
                        net = r(slim.fully_connected(net, 256, scope='6_fc'))
                        #net = r(slim.dropout(net, scope='6_dropout'))
                        net = r(slim.fully_connected(net, 128, scope='7_fc'))

                    encoded = net

                    dstack = []
                    def r(x):
                        dstack.append(x)
                        return x
                    with tf.variable_scope("decoder"):
                        net = r(slim.fully_connected(net, 256, scope='1_fc'))
                        #net = r(slim.dropout(net, scope='1_dropout'))
                        net = r(slim.fully_connected(net, 4*4*48, scope='2_fc'))
                        net = r(tf.reshape(net, [-1, 4, 4, 48], name='2_reshape'))
                        net = r(upscale(net, [2, 2], name='3_upscale'))
                        net = r(slim.conv2d_transpose(net, 42, [4, 4], scope='4_conv'))
                        net = r(upscale(net, [2, 2], name='5_upscale'))
                        net = r(slim.conv2d_transpose(net, 36, [3, 3], scope='6_conv'))

```

```
ope='6_conv'))  
net = r(slim.conv2d_transpose(net, 3, [4, 4], scope='7_conv'))  
  
reconstructed = net  
  
return (encoded, reconstructed, estack, dstack)
```

```

def autoencoder_model_3(image_batch):
    with slim.arg_scope([slim.conv2d, slim.fully_connected],
                        activation_fn=tf.nn.relu,
                        weights_initializer=tf.truncated_normal_initializer(stddev=0.01),
                        biases_initializer=tf.zeros_initializer,
                        #weights_regularizer=slim.l2_regularizer(0.0005),
                        biases_regularizer=None):
        with slim.arg_scope([slim.conv2d, slim.conv2d_transpose],
                            padding='VALID'):
            with slim.arg_scope([slim.dropout], keep_prob=0.8):
                with slim.arg_scope([slim.conv2d, slim.fully_connected, slim.conv2d_transpose],
                                    normalizer_fn=slim.batch_norm):
                    net = image_batch

                    estack = []
                    def r(x):
                        estack.append(x)
                        return x
                    with tf.variable_scope("encoder"):
                        net = r(slim.conv2d(net, 36, [4, 4], scope='1_conv'))
                        net = r(slim.conv2d(net, 42, [3, 3], scope='2_conv'))
                        net = r(slim.max_pool2d(net, [2, 2], scope='3_max_pool'))
                        net = r(slim.conv2d(net, 48, [4, 4], scope='4_conv'))
                        net = r(slim.max_pool2d(net, [2, 2], scope='5_max_pool'))
                        net = r(slim.flatten(net, scope='5_flatten'))
                        net = r(slim.fully_connected(net, 256, scope='6_fc'))
                        #net = r(slim.dropout(net, scope='6_dropout'))
                        net = r(slim.fully_connected(net, 128, scope='7_fc'))

                    encoded = net

                    dstack = []
                    def r(x):
                        dstack.append(x)
                        return x
                    with tf.variable_scope("decoder"):
                        net = r(slim.fully_connected(net, 256, scope='1_fc'))
                        #net = r(slim.dropout(net, scope='1_dropout'))
                        net = r(slim.fully_connected(net, 4*4*48, scope='2_fc'))
                        net = r(tf.reshape(net, [-1, 4, 4, 48], name='2_reshape'))
                        net = r(upscale(net, [2, 2], name='3_upscale'))
                        net = r(slim.conv2d_transpose(net, 42, [4, 4], scope='4_conv'))
                        net = r(upscale(net, [2, 2], name='5_upscale'))
                        net = r(slim.conv2d_transpose(net, 36, [3, 3], scope='6_conv'))

```



```

ope='6_conv'))
        net = r(slim.conv2d_transpose(net, 3, [4, 4], activation_fn=tf.nn.sigmoid, scope='7_conv'))

        reconstructed = net

        return (encoded, reconstructed, estack, dstack)

```

```

In [16]: if sess is not None:
        sess.close()
        tf.reset_default_graph()
        sess = tf.InteractiveSession()
        patch_tensor = tf.placeholder(dtype='float32', shape=(None, 27, 27, 3))
        with tf.variable_scope("autoencoder"):
            with slim.arg_scope([slim.dropout], is_training=True), slim.arg_scope(
                [slim.conv2d, slim.conv2d_transpose, slim.fully_connected], normalizer_params={
                    'is_training': True}):
                (encoded_tensor, reconstructed_tensor, estack, dstack) = autoencoder_model_3(patch_tensor)
            with tf.variable_scope("autoencoder", reuse=True):
                with slim.arg_scope([slim.dropout], is_training=False), slim.arg_scope(
                    [slim.conv2d, slim.conv2d_transpose, slim.fully_connected], normalizer_params={
                        'is_training': False}):
                    (eval_encoded_tensor, eval_reconstructed_tensor, eval_estack, eval_dstack) = autoencoder_model_3(patch_tensor)

        print "Encoded:", encoded_tensor.get_shape().as_list()
        print "Reconstructed:", reconstructed_tensor.get_shape().as_list()

```

```

Encoded: [None, 128]
Reconstructed: [None, 27, 27, 3]

```

```

In [17]: # Restore the (final) model
        saver = tf.train.Saver()
        saver.restore(sess, "autoencoder_models/final_model_v1/v150/model.ckpt")

```

```

In [19]: def get_training_op(reconstructed_tensor, patch_tensor, lr_tensor):
        difference = reconstructed_tensor - patch_tensor
        #print difference.get_shape().as_list()
        square_difference = tf.square(difference)
        mean_squared_error = tf.reduce_mean(square_difference)
        slim.losses.add_loss(mean_squared_error)
        total_loss = slim.losses.get_total_loss()
        optimizer = tf.train.AdamOptimizer(learning_rate=lr_tensor)#learning_rate=0.01
        #optimizer = tf.train.AdamOptimizer()
        train_op = slim.learning.create_train_op(total_loss, optimizer)
        return (train_op, mean_squared_error)

```

```

In [20]: lr_tensor = tf.placeholder(dtype='float32', shape=[])
        (train_op, loss) = get_training_op(reconstructed_tensor, patch_tensor, lr_tensor)

```

```

def train_loop(sess, patch_tensor, reconstructed_tensor, eval_reconstruct
ed_tensor, train_patches, test_patches, train_op, loss_tensor, epochs, ba
tch_size, reset=True):
    if not os.path.exists("imgs/train"):
        os.makedirs("imgs/train")
    if not os.path.exists("imgs/test"):
        os.makedirs("imgs/test")
    tr_loss = []
    tst_loss = []
    N = train_patches.shape[0]
    if reset:
        sess.run(tf.initialize_all_variables())
    for e in xrange(epochs):
        for i in xrange(0, N, batch_size):
            [total_loss, loss] = sess.run([train_op, loss_tensor], feed_d
ict={
                lr_tensor:0.001,
                patch_tensor:train_patches[i:i+batch_size],
            })
            step = i / batch_size
            if step % 50 == 0:
                [test_loss] = sess.run([loss_tensor], feed_dict={
                    patch_tensor:test_patches[:100],
                })
                test_loss = 0
                print "Epoch %d, step %d, total loss %f, training loss
%f, test_loss %f" % (e, step, total_loss, loss, test_loss)
                tr_loss.append(loss)
                tst_loss.append(test_loss)
                #print "Starting print"
                imgs = sess.run(eval_reconstructed_tensor, feed_dict={
                    patch_tensor:train_patches[:16],
                })
                #plt.figure()
                for i in range(16):
                    plt.subplot(4,4,i+1)
                    plt.imshow(imgs[i])
                plt.savefig("imgs/train/%d_%d.png" % (e, step))
                #print "Done print training"
                imgs = sess.run(eval_reconstructed_tensor, feed_dict={
                    patch_tensor:test_patches[:16],
                })
                #plt.figure()
                for i in range(16):
                    plt.subplot(4,4,i+1)
                    plt.imshow(imgs[i])
                plt.savefig("imgs/test/%d_%d.png" % (e, step))
                #print "Done print testing"
            # End-of-epoch printing
            [test_loss] = sess.run([loss_tensor], feed_dict={
                patch_tensor:test_patches[:100],
            })
            test_loss = 0
            print "End of epoch %d, training loss %f, test_loss %f" % (e, los
s, test_loss)
            # Save the model
            saver = tf.train.Saver()

```

```
os.makedirs("autoencoder_models/autosave/v%d" % e)
save_path = saver.save(sess, "autoencoder_models/autosave/v%d/model.ckpt" % e)
print "Saved to:", save_path
#print "Final training loss %f" % (loss)
return (tr_loss, tst_loss)
```

```
In [ ]: tr_loss, tst_loss = train_loop(  
        sess, patch_tensor, reconstructed_tensor, eval_reconstructed_tensor,  
        train_dict['patches'], test_patches, train_op, loss, 500, 100, reset=False)  
e)
```


```
Epoch 35, step 50, total loss 0.004751, training loss 0.004751, test_loss 0.000000
Epoch 35, step 100, total loss 0.004825, training loss 0.004825, test_loss 0.000000
Epoch 35, step 150, total loss 0.004640, training loss 0.004640, test_loss 0.000000
Epoch 35, step 200, total loss 0.004874, training loss 0.004874, test_loss 0.000000
Epoch 35, step 250, total loss 0.004375, training loss 0.004375, test_loss 0.000000
```

```
In [332]: # Save the model
saver = tf.train.Saver()
save_path = saver.save(sess, "autoencoder_models/v3/model.ckpt")
print "Saved to:", save_path
```

Saved to: autoencoder\_models/v3/model.ckpt

## Compare Reconstructions to Originals

```
In [325]: for i in range(16):
plt.subplot(4,4,i+1)
plt.imshow(train_dict['patches'][i])
plt.savefig("train_golden.png")
for i in range(16):
plt.subplot(4,4,i+1)
plt.imshow(test_patches[i])
plt.savefig("test_golden.png")
```



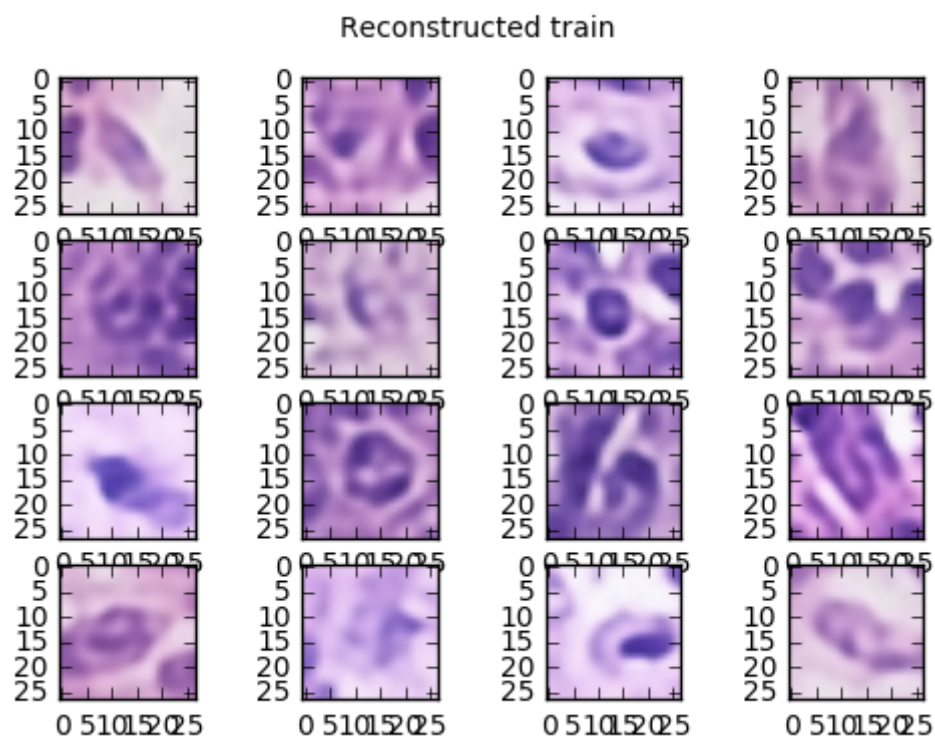
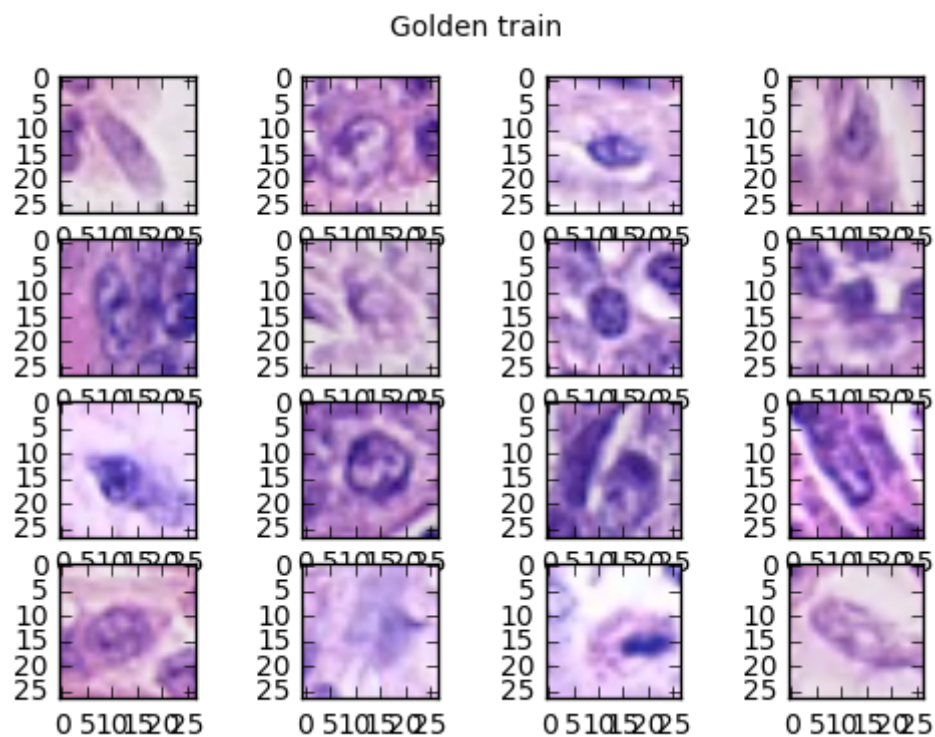
```
In [26]: # Compare golden to best recreations
plt.figure()
for i in range(16):
    plt.subplot(4,4,i+1)
    plt.imshow(train_dict['patches'][i])
plt.suptitle("Golden train")

plt.figure()
imgs = sess.run(eval_reconstructed_tensor, feed_dict={
    patch_tensor:train_dict['patches'][:16],
})
for i in range(16):
    plt.subplot(4,4,i+1)
    plt.imshow(imgs[i])
plt.suptitle("Reconstructed train")


plt.figure()
for i in range(16):
    plt.subplot(4,4,i+1)
    plt.imshow(test_patches[i])
plt.suptitle("Golden test")

plt.figure()
imgs = sess.run(eval_reconstructed_tensor, feed_dict={
    patch_tensor:test_patches[:16],
})
for i in range(16):
    plt.subplot(4,4,i+1)
    plt.imshow(imgs[i])
plt.suptitle("Reconstructed test")
```


Out[26]: <matplotlib.text.Text at 0x7fd7f73e1b90>



Golden test



Reconstructed test



## Save Encodings



```
In [44]: # Calculate the encodings for everything in the training and test sets
train_N = train_dict['patches'].shape[0]
test_N = test_patches.shape[0]
features_N = eval_encoded_tensor.get_shape().as_list()[1]
train_encodings = np.zeros((train_N, features_N))
test_encodings = np.zeros((test_N, features_N))
batch_size = 100 # We will do it in batches of 100 to save on memory

# First do the train set
for i in xrange(0, train_N, batch_size):
    train_encodings[i:i+batch_size] = sess.run(eval_encoded_tensor, feed_dict={
        patch_tensor: train_dict['patches'][i:i+batch_size]
    })

# Next do the test set
for i in xrange(0, test_N, batch_size):
    test_encodings[i:i+batch_size] = sess.run(eval_encoded_tensor, feed_dict={
        patch_tensor: test_patches[i:i+batch_size]
    })
```

```
In [45]: # Okay now we want to write these out to disk
np.save("train_autoencodings.npy", train_encodings)
np.save("test_autoencodings.npy", test_encodings)
```

## Feature-Space Manipulations

```
In [19]: categories
```

```
Out[19]: ['epithelial', 'fibroblast', 'inflammatory', 'others']
```