

```
In [1]: %matplotlib inline
```

```
In [2]: import numpy as np
import matplotlib
import matplotlib.pyplot as plt
```

```
In [3]: import classifier_utils as utils
```

```
In [4]: import cv2
```

```
In [5]: np.seterr(divide='ignore')
```

```
Out[5]: {'divide': 'warn', 'invalid': 'warn', 'over': 'warn', 'under': 'ignore'}
```

Import SVM Implementation

```
In [173]: # TODO: shouldn't polute the namespace like this, svm is a pretty short p
refix!
from svm import get_gradient
from svm import get_activations
from svm import get_cost
from svm import get_cost_and_gradient
from svm import training_step
from svm import get_multiclass_predictions
from svm import train_loop
from svm import get_np_weighted_f1
from svm import status_report
from svm import multiclass_train_loop
```

Simple Dataset SVM Example


```
In [164]: data = np.array([[ -4, 10], # class 0
                           [ 1, 10],
                           [-6, 8],
                           [-3, 8],
                           [ 3, 8],
                           [-1, 7],
                           [-4, 5],
                           [ 0, 5],
                           [-6, 3],
                           [-2, 2],
                           [-6, 1],
                           [-5, -1],
                           [-7, -2],
                           [-6, -4],

                           [-1, -2], # class 1
                           [-4, -5],
                           [-1, -5],
                           [ 1, -5],
                           [-6, -7],
                           [-2, -7],
                           [ 1, -7],
                           [ 3, -7],
                           [-4, -8],
                           [ 1, -9],
                           [ 5, -9],
                           [-7, -10],
                           [-4, -10],
                           [-2, -10],

                           [ 7, 10], # class 2
                           [ 6, 7],
                           [ 7, 6],
                           [ 3, 5],
                           [ 5, 4],
                           [ 7, 3],
                           [ 2, 2],
                           [ 5, 2],
                           [ 5, 0],
                           [ 2, -1],
                           [ 7, -1],
                           [ 5, -2],
                           [ 5, -4],
                           [ 7, -6]])
labels = np.array([[1, 0, 0]]*14 + [[0, 1, 0]]*14 + [[0, 0, 1]]*14)
```

```
In [165]: # Plot dataset  
plt.scatter(data[:14,0], data[:14,1], c='r')  
plt.scatter(data[14:28,0], data[14:28,1], c='g')  
plt.scatter(data[28:,0], data[28:,1], c='b')
```

Out[165]: <matplotlib.collections.PathCollection at 0x7f5369b73690>



```
In [172]: # Run SVM
C = 10
learning_rate = 0.003
batch_size = 2800 # Just some really big number to mean "the entire dataset in one batch"
epochs = 1000
ws = multiclass_train_loop(data, labels, data, labels, data, labels, C, learning_rate, batch_size, epochs, quiet=True)
print ws
```

Epoch FINAL, step FINAL:

```
    train loss [1.0913469653460335, 2.5954427415520955, 2.5829900778049462],
    validation loss [1.0913469653460335, 2.5954427415520955, 2.5829900778049462],
    test loss [1.0913469653460335, 2.5954427415520955, 2.5829900778049462],
    per-category test accuracy [ 1.          0.90476191  0.90476191],
    overall validation accuracy 1.000000,
    overall test accuracy 1.000000,
    validation f1 1.000000,
    test f1 1.000000
[[-0.55644254 -0.0645514  0.39326842]
 [ 0.4173721  -0.29224589  0.06284798]
 [-0.72548389 -0.8712765  -0.83905256]]
```

```
In [174]: # Visualize results
# Dots are correctly classifier, crosses are incorrectly classifier
# Coloured lines are decision boundaries for each classifier
(predictions, _) = get_multiclass_predictions(data, ws)
actual = np.argmax(labels, axis=1)
correct = np.equal(predictions, actual)

mask = np.logical_and(correct, np.equal(actual, 0))
plt.scatter(data[mask, 0], data[mask,1], c='r')
mask = np.logical_and(np.logical_not(correct), np.equal(actual, 0))
plt.scatter(data[mask, 0], data[mask,1], c='r', marker='+')

mask = np.logical_and(correct, np.equal(actual, 1))
plt.scatter(data[mask, 0], data[mask,1], c='g')
mask = np.logical_and(np.logical_not(correct), np.equal(actual, 1))
plt.scatter(data[mask, 0], data[mask,1], c='g', marker='+')

mask = np.logical_and(correct, np.equal(actual, 2))
plt.scatter(data[mask, 0], data[mask,1], c='b')
mask = np.logical_and(np.logical_not(correct), np.equal(actual, 2))
plt.scatter(data[mask, 0], data[mask,1], c='b', marker='+')


plt.axis((-8, 8, -15, 15))

a, b, c = ws[:, 0]
x = np.linspace(-8, 8)
y = -a/b * x - c/b
plt.plot(x, y, c='r')

a, b, c = ws[:, 1]
x = np.linspace(-8, 8)
y = -a/b * x - c/b
plt.plot(x, y, c='g')

a, b, c = ws[:, 2]
x = np.linspace(-8, 8)
y = -a/b * x - c/b
plt.plot(x, y, c='b')
```

Out[174]: [`<matplotlib.lines.Line2D at 0x7f537f2d46d0>`]



Get Dataset

```
In [6]: # To save memory, delete old dataset before creating new one
try:
    del test, train, validation # if they exist
except:
    pass

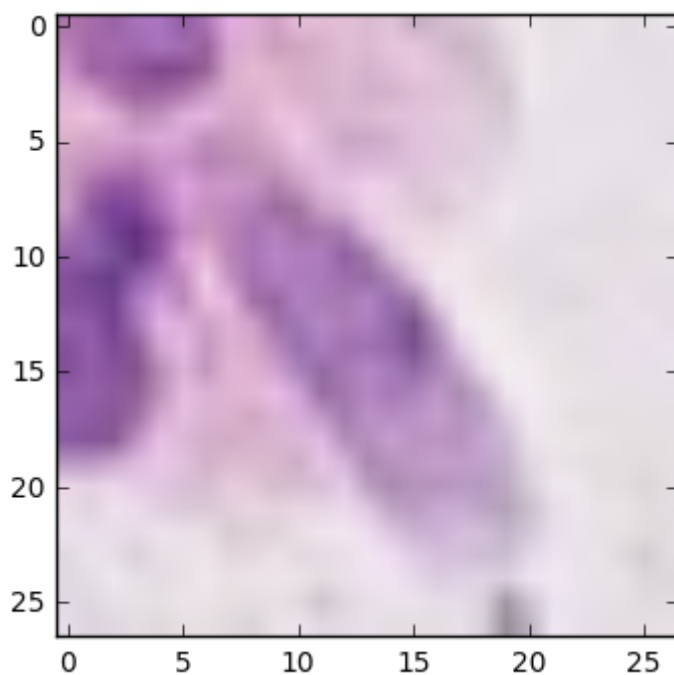
# Dataset
categories = [
    'epithelial',
    'fibroblast',
    'inflammatory',
    'others',
]
train, test = utils.get_augmented_dataset(categories)

# Carve out a validation set from our test set
# Split it 50/50
validation = {}
for k in list(test.iterkeys()):
    values = test[k]
    N = len(values)
    new_N = N/2
    test[k], validation[k] = np.split(values, [new_N])
```

Dropped 2082 patches because too close to image border

```
In [7]: plt.imshow(train['patches'][0])
```

```
Out[7]: <matplotlib.image.AxesImage at 0x7f537f0e3090>
```



Get HOG Features

```
In [8]: def get_hog_features(hog, patches):  
        N = patches.shape[0]  
        descriptor_size = hog.getDescriptorSize()  
        results = np.zeros((N, descriptor_size))  
        for (i, patch) in enumerate(patches):  
            uint_patch = (patch * 255).astype('uint8')  
            resized_patch = cv2.resize(uint_patch, winSize, interpolation=cv2.  
TER_LINEAR)  
            feature_vector = hog.compute(resized_patch)[: , 0]  
            results[i] = feature_vector  
        return results
```



```
In [9]: # Compute HOG descriptors for each test/train example

# Specify HOG parameters
winSize = (32, 32)
blockSize = (16, 16)
blockStride = (8, 8)
cellSize = (8, 8)
nbins = 9
hog = cv2.HOGDescriptor(winSize, blockSize, blockStride, cellSize, nbins)

# Get training descriptors
train['hog_features'] = get_hog_features(hog, train['patches'])

# Get validation descriptors
validation['hog_features'] = get_hog_features(hog, validation['patches'])

# Get test descriptors
test['hog_features'] = get_hog_features(hog, test['patches'])
```

Load Autoencoder Features

```
In [12]: train['autoencodings'] = np.load("train_autoencodings.npy")
test_encodings = np.load("test_autoencodings.npy")

# Split test encodings into validation and test
new_N = test['patches'].shape[0]
test['autoencodings'], validation['autoencodings'] = np.split(test_encodings, [new_N])
```

```
In [13]: for (k, v) in train.iteritems():
          print "train", k, v.shape
          for (k, v) in validation.iteritems():
              print "validation", k, v.shape
          for (k, v) in test.iteritems():
              print "test", k, v.shape
```

```
train hog_features (60000, 324)
train hsv_factors (60000, 3)
train deltas (60000, 2)
train patches (60000, 27, 27, 3)
train rots (60000,)
train labels (60000, 4)
train flips (60000,)
train centres (60000, 2)
train autoencodings (60000, 128)
train img_ids (60000,)
validation hog_features (2037, 324)
validation patches (2037, 27, 27, 3)
validation labels (2037, 4)
validation centres (2037, 2)
validation autoencodings (2037, 128)
validation img_ids (2037,)
test hog_features (2036, 324)
test patches (2036, 27, 27, 3)
test labels (2036, 4)
test centres (2036, 2)
test autoencodings (2036, 128)
test img_ids (2036,)
```

Functions for Model Selection

```

def output_global_status(C, prefix):
    # Helper function to access per-category time series
    def g(series, category):
        return [x[category] for x in series]

    def save_plot(suffix):
        full_dir = os.path.join(prefix, suffix)
        full_name = os.path.join(full_dir, "C=%010.4f.png" % C)
        if not os.path.exists(full_dir):
            os.makedirs(full_dir)
        plt.savefig(full_name)
        plt.clf()

    best_by_f1 = np.argmax(val_f1)
    print "Best according to f1: index %d, val f1 %f, test f1 %f" % (
        best_by_f1, val_f1[best_by_f1], test_f1[best_by_f1])
    best_by_acc = np.argmax(val_accuracy)
    print "Best according to accuracy: index %d, val acc %f, test acc %f"
    % (
        best_by_acc, val_accuracy[best_by_acc],
        test_accuracy[best_by_acc])

    x = range(len(g(loss, 0)))
    plt.plot(x, g(loss, 0), 'r')
    plt.plot(x, g(loss, 1), 'g')
    plt.plot(x, g(loss, 2), 'b')
    plt.plot(x, g(loss, 3), 'y')
    plt.title("C = %f, train loss" % C)
    save_plot("train_loss")
    x = range(len(g(test_loss, 0)))
    plt.plot(x, g(test_loss, 0), 'r')
    plt.plot(x, g(test_loss, 1), 'g')
    plt.plot(x, g(test_loss, 2), 'b')
    plt.plot(x, g(test_loss, 3), 'y')
    plt.title("C = %f, test loss" % C)
    save_plot("test_loss")
    x = range(len(g(val_loss, 0)))
    plt.plot(x, g(val_loss, 0), 'r')
    plt.plot(x, g(val_loss, 1), 'g')
    plt.plot(x, g(val_loss, 2), 'b')
    plt.plot(x, g(val_loss, 3), 'y')
    plt.title("C = %f, val loss" % C)
    save_plot("val_loss")

    plt.plot(x, val_accuracy)
    plt.title("C = %f, overall val accuracy" % C)
    save_plot("overall_val_accuracy")

    plt.plot(x, test_accuracy)
    plt.title("C = %f, overall test accuracy" % C)
    save_plot("overall_test_accuracy")

    plt.plot(x, val_f1)
    plt.title("C = %f, overall val f1" % C)
    save_plot("overall_val_f1")

    plt.plot(x, test_f1)

```

```
plt.title("C = %f, overall test f1" % C)
save_plot("overall_test_f1")

return (best_by_f1, best_by_acc)
```

```
In [108]: def svm_C_search(train, validation, test, feature_name, learning_rate, batch_size, epochs, Cs, prefix):
    results = {
        'Cs' : [],
        'best_validation_f1_index': [],
        'best_validation_accuracy_index': [],
        'best_ws_by_validation_f1': [],
        'best_ws_by_validation_accuracy': [],
        'best_validation_f1': [],
        'best_validation_accuracy': [],
        'test_f1_at_best': [],
        'test_accuracy_at_best': [],
    }

    # Run SVM with various values of C
    for C in Cs:
        print "C=%f" % C
        multiclass_train_loop(train[feature_name], train['labels'],
                               validation[feature_name], validation['labels'],
                               test[feature_name], test['labels'],
                               C, learning_rate, batch_size, epochs,
                               quiet=True)
        (f1_index, accuracy_index) = output_global_status(C, prefix)
        results['Cs'].append(C)
        results['best_validation_f1_index'].append(f1_index)
        results['best_validation_accuracy_index'].append(accuracy_index)
        results['best_ws_by_validation_f1'].append(params[f1_index])
        results['best_ws_by_validation_accuracy'].append(params[accuracy_index])
        results['best_validation_f1'].append(val_f1[f1_index])
        results['best_validation_accuracy'].append(val_accuracy[accuracy_index])
        results['test_f1_at_best'].append(test_f1[f1_index])
        results['test_accuracy_at_best'].append(test_accuracy[accuracy_index])

        print ""
        print ""

    # Print a summary message and return the results
    best_f1_C_index = np.argmax(results['best_validation_f1'])
    best_accuracy_C_index =
np.argmax(results['best_validation_accuracy'])
    print "OVERALL"
    print "Best f1 is at index %d, best accuracy is at index %d" % (best_f1_C_index, best_accuracy_C_index)

    return results
```

Apply SVM to HOG Features

```
In [109]: # Run SVM with various values of C
learning_rate = 0.0003
batch_size = 200
epochs = 100
# Recommended to try C in range  $2^{-5} \dots 15$  from http://www.csie.ntu.edu.tw/~cjlin/papers/guide/guide.pdf
Cs = [2.0**x for x in xrange(-5, 16)]
hog_results = svm_C_search(
    train, validation, test, 'hog_features', learning_rate, batch_size, epochs, Cs, 'svm_graphs/hog')
```

OVERALL

Best f1 is at index 16, best accuracy is at index 16

<matplotlib.figure.Figure at 0x7f53784cd4d0>

```
In [154]: # Store results
# import pickle
# with open('hog_svn_results.p', 'wb') as f:
#     pickle.dump(hog_results, f)
# np.save('hog_svn_best_ws.npy', np.hstack(hog_results['best_ws_by_validation_f1'][16]))
```

```
In [157]: # Restore results
import pickle
with open('hog_svn_results.p', 'rb') as f:
    hog_results = pickle.load(f)
```


```
In [160]: print hog_results['test_accuracy_at_best'][16]
print hog_results['test_f1_at_best'][16]
```

0.556483
0.526957574876


```
In [162]: hog_ws = np.hstack(hog_results['best_ws_by_validation_f1'][16])
hog_predictions, _ = get_multiclass_predictions(test['hog_features'], hog_ws)
hog_actuals = np.argmax(test['labels'], axis=1)
print np.count_nonzero(np.equal(hog_predictions, hog_actuals)) / float(len(hog_actuals))
print get_np_weighted_f1(hog_predictions, hog_actuals, 4)
```

0.548133595285
0.526957574876

```
In [120]: plt.plot(hog_results['Cs'], hog_results['best_validation_accuracy'])  
plt.plot(hog_results['Cs'], hog_results['test_accuracy_at_best'])  
plt.xscale('log')
```



```
In [116]: plt.plot(hog_results['Cs'], hog_results['best_validation_f1'])  
plt.plot(hog_results['Cs'], hog_results['test_f1_at_best'])  
plt.xscale('log')
```




```
In [164]: (predictions, _) = get_multiclass_predictions(validation['hog_features'],  
ws)  
print np.count_nonzero(np.equal(predictions, 0))  
print np.count_nonzero(np.equal(predictions, 1))  
print np.count_nonzero(np.equal(predictions, 2))  
print np.count_nonzero(np.equal(predictions, 3))
```

832

921

284

0

Apply SVM to Autoencoder Features

```
In [121]: # Run SVM with various values of C
learning_rate = 0.0001
batch_size = 200
epochs = 200 # Many of these take longer to converge than HOG, so run for
             200 epochs
# Recommended to try C in range  $2^{-5} \dots 15$  from http://www.csie.ntu.edu
# u.tw/~cjlin/papers/guide/guide.pdf
Cs = [2.0**x for x in xrange(-5, 16)]
auto_results = svm_C_search(
    train, validation, test, 'autoencodings', learning_rate, batch_size,
    epochs, Cs, 'svm_graphs/auto')
```

OVERALL

Best f1 is at index 10, best accuracy is at index 10

<matplotlib.figure.Figure at 0x7f537ef7a250>

```
In [146]: # Store results
#import pickle
#with open('auto_svn_results.p', 'wb') as f:
#    pickle.dump(auto_results, f)
#np.save('auto_svn_best_ws.npy', np.hstack(auto_results['best_ws_by_validation_f1'][10]))
```

```
In [147]: # Restore results
import pickle
with open('auto_svn_results.p', 'rb') as f:
    auto_results = pickle.load(f)
```


```
In [149]: print auto_results['test_accuracy_at_best'][10]
print auto_results['test_f1_at_best'][10]
```

0.663065
0.657819415973

```
In [153]: auto_ws = np.hstack(auto_results['best_ws_by_validation_f1'][10])
auto_predictions, _ = get_multiclass_predictions(test['autoencodings'], auto_ws)
#print auto_predictions[:100]
auto_actuals = np.argmax(test['labels'], axis=1)
#print auto_actuals[:100]
print np.count_nonzero(np.equal(auto_predictions, auto_actuals)) /
float(len(auto_actuals))
print get_np_weighted_f1(auto_predictions, auto_actuals, 4)
```

0.669941060904
0.657819415973

```
In [122]: plt.plot(auto_results['Cs'], auto_results['best_validation_accuracy'])  
plt.plot(auto_results['Cs'], auto_results['test_accuracy_at_best'])  
plt.xscale('log')
```



```
In [125]: plt.plot(auto_results['Cs'], auto_results['best_validation_f1'])  
plt.plot(auto_results['Cs'], auto_results['test_f1_at_best'])  
plt.xscale('log')
```

